

What a reading proposal should state in order to be assessable: a minimum protocol for undeciphered scripts

Alberto Acedo Biome Makers Inc. Draft v0.1, 9 September 2026. Intended for a methodological opinion note (Antiquity, Journal of Archaeological Science: Reports, or the forum of a field journal). Not to be submitted before the empirical papers have a destination.

Abstract

Proposals for reading undeciphered scripts appear faster than the field can evaluate them, and the reason is not a shortage of expertise but the absence of a shared minimum of what a proposal must state to be assessable at all. Drawing on the experience of applying a single procedure to five corpora with known or partly known answers (Etruscan, Eteocypriot, Proto-Elamite, proto-cuneiform and Iberian) and to one that has none (Linear A), and on our own withdrawn results, this paper proposes eight requirements: a null model for every distributional claim; a positive control in a corpus with a known answer; a prediction about unseen text; a declaration of how many tests were run, with correction for multiple comparisons; verification of the data in the primary edition; publication of negatives with their figures; an account of the ending of any sign-group identified with a word of another language, or an explicit statement that it is unaccounted for; and, before any absence is interpreted, a check that the instrument could have seen what is said to be missing. None of the eight requires computational methods, and none demands that a proposal be correct: they demand that it be evaluable. For each requirement a case is given in which its absence produced a false positive, and four of those cases are the author's own.

Keywords: undeciphered scripts, method, null models, reproducibility, Linear A, Indus script.

1. The problem

Reading proposals appear faster than the field can evaluate them, and the million-dollar prize for the Indus script (2025) has accelerated the pace. The problem is not a lack of competence: it is that there is no common minimum of what a proposal must declare for another to judge it. Without that minimum, discussion is settled by authority or by exhaustion, and a proposal that is wrong for a demonstrable reason survives beside one that is wrong for reasons nobody can articulate.

This paper proposes no method and evaluates no one else's proposal. It proposes eight requirements, all elementary in other empirical sciences, and for each shows a case in which its absence produced a false positive. Four of the cases are the author's own, and they are there deliberately: whoever proposes a protocol should first show where he has failed it.

2. The eight requirements

2.1 A null model for every distributional claim

Every claim of the type “this word occurs with that one”, “this affix is productive”, “these two document classes differ” needs a permutation that preserves what the claim does not purport to explain and destroys only what it does. Any null will not do: the wrong null produces positives.

Own case: we measured the ascendancy-to-capacity ratio of the document-word network at Haghia Triada and at Susa, and both fell within Ulanowicz’s “window of vitality”, an attractive result. With a null preserving row and column sums, the null reproduces the observed value to three decimal places. The index was measuring the marginal distributions. Withdrawn.

Own case 2: we measured the compositionality of the Minoan lexicon (whether long words decompose into recurrent elements) and obtained 0.12 against 0.02 for the shuffled-syllable null, above Greek at equal size. With a null preserving syllable transitions, the observed value falls within the null ($p = 0.60$): it was the phonotactics of the syllabary. Withdrawn.

Rule: the null must preserve everything that is not the hypothesis.

2.2 A positive control in a corpus with a known answer

Before applying an instrument to the unknown, one must check that it recovers the known in a comparable corpus. If it does not, the instrument’s silence is not information.

Our use: we calibrated on Etruscan (it recovers the genitive before *clan*, the chronology of syncope and the geography of the sibilants), on Eteocypriot against Cypriot Greek (it separates the two languages by their endings), on Proto-Elamite (it recovers Dahl’s classification of the numerical systems, the institutional heading and the name/commodity partition) and on Iberian (it recovers the formulaic elements).

Counterexample from the field: most reading proposals for Linear A and for the Indus script report no positive control, so it is not known whether their method would see the answer if it were in front of them.

2.3 A prediction about unseen text

A reading that only explains the texts already known cannot be evaluated. The proposal must say what it expects to find in the next inscription and what would refute it.

Our use: the reading of the KU-PA family predicts that its members will go on appearing with the cyperus logogram and with dry commodities, that a new roundel with KA-KU-PA will carry that logogram or a ligature of it, and that the family will not appear as the recipient of liquids. The protocol for what will be done with the Knossos sceptre when it is published is pre-registered before the text is known.

Note: the prediction does not require that the text appear; it requires that it be written down beforehand.

2.4 How many tests were run, with correction for multiple comparisons

A procedure that tests forty hypotheses finds two below 0.05 by chance. The proposal must declare the denominator and correct within each family of tests.

Own case: of eleven suffixes tested in Linear A, three exceed the threshold uncorrected and none survives Bonferroni; we report it as such, and the result we sustain is not “these three suffixes are productive” but “there is a positional signal of morphology that none of the suffixes individually demonstrates”.

Counterexample from the field: screens for lexical coincidence between an unread script and an extensive dictionary, which are in effect thousands of simultaneous tests, almost never declare that number.

2.5 Verify the datum in the primary edition

Derived transliterations, including excellent ones, introduce errors that propagate.

Own case: one of our supports for reading KU-PA as a commodity was a ligature, *304+PA-KU-PA, on a Haghia Triada nodule. The palaeographic database reads three consecutive signs there: the ligature does not exist. It sustained an argument for two days.

Own case 2: the same source transcribes some fraction signs as whole numbers, which made the arithmetic of several tablets fail.

Rule: every datum sustaining a claim is checked in the edition or in the palaeographic database, and the version used is declared.

2.6 Publish negatives with their figures

A negative without a number is not information: “we found no co-exclusion” may mean that there is none or that the corpus is small. With a number, the negative is a measure of power and saves the next investigator the work.

Our use: there is no measurable co-exclusion between Linear A words (zero pairs with expectation above one in 224 documents), and there is in Proto-Elamite with 830 tablets (13 pairs); a network descriptor has no test-retest reliability at these sizes; comparison by sound does not discriminate between candidate languages (all between 1.3 and 1.7 over their own null) either in Linear A or in Linear B, where the same rule produces 383 false positives.

2.7 Account for the ending, or state that it is unaccounted for

When a proposal identifies a sign-group with a known word of another language, the comparison usually rests on the stem and leaves the ending untreated. That is where most false positives enter, because a stem of two or three syllables coincides easily and an unexplained ending may be concealing morphology, a particle, or a misplaced word boundary.

The requirement is cheap: **state what the ending is proposed to be, or state expressly that it is unaccounted for.** It does not demand being right; it demands not being silent.

An own case illustrates it, which is why it is included here. In the companion paper on names formed on plant words we proposed KU-MI-NA-QE as formed on the word for cumin (Mycenaean ku-mi-no, Greek κύμινον, a Semitic loan), on the basis of the shared stem and of its distribution on a sealed roundel, where the sign-group is the party and the logogram the commodity. The vowel -a against Mycenaean -o is no objection, since it follows the correspondence rule measured over 46 pairs (the Minoan final vowel becomes -o in Greek and -a is never preserved). **But the ending -QE was left untreated**, and a consultation with J. M. Jiménez Delgado (University of Seville) identified it as the problem with the case.

Checked in the corpus, the objection holds. There are seven sign-groups ending in -qe among 782 types (0.9%), and only two have their stem attested separately; against the shuffled-syllable null preserving lengths, those two pairs give $p = 0.086$, so -qe does not hold as a productive suffix, unlike -ja, which does with seven pairs. And KA-PA-QE is telling: its stem KA-PA is a Cretan toponym (Uchitel 2002-2003) heading six Haghia Triada tablets, so that KA-PA-QE reads naturally as “and KA-PA”. In Mycenaean, -qe is the copulative enclitic. If Minoan had a comparable enclitic, KU-MI-NA-QE would be the word ku-mi-na with a particle attached rather than a name, and its two contexts admit that reading.

The case is not refuted; it is left **open**: it satisfies two of the three conditions, with the third in dispute. Writing it that way is what this requirement asks, and it is what the original proposal did not do.

2.8 Before interpreting an absence, check that the instrument could have seen it

A zero is the most tempting observation in a small corpus and the easiest to misread. When something does not appear, there are two possible explanations and only one is interesting: that it is absent because the object did not have it, or that it is absent because the instrument could not see it. The requirement consists in ruling out the second before asserting the first.

Two own examples illustrate it, both detected on applying this requirement retroactively.

First. We claimed that no Linear A sign-group ends in -so, and contrasted this with the fact that Cretan toponyms end in -so in Greek (Knossos, Tyliossos, Amnisos). Measuring the full distribution of syllables, it turns out that **the syllable so never occurs in the corpus, in any position**, and with it seven others of the -o column (do, jo, mo, no, qo, wo, zo). The word-final absence was not a phonotactic feature of Minoan: it was a gap in the transliterated syllabary. The claim was withdrawn and replaced by the comparison of each syllable with its own general frequency, which does measure a positional preference and does not depend on which signs have an assigned value.

Second. Applying by semantic field the functional criterion of a companion paper, names formed on animal words gave zero matches against nine for the plant field. The immediate reading was that Minoan onomastics did not form names on animals, which would be a cultural feature with content. But the search is necessarily made with Mycenaean stems, and in Crete livestock is always written with a logogram and never with syllables: the Minoan words for

animals with which to search do not exist. The zero measured the absence of available vocabulary, not the absence of the pattern.

The check is cheap in both cases and consists of the same question: what would have to be in the corpus for the instrument to detect what I say is missing? If that prior condition is not met, the absence is not a datum about the language.

The requirement is therefore stated thus: **every claim of the form “X does not occur” must be accompanied by the check that X could have occurred**, that is, that the sign, word or category was available to the instrument in the material and the transliteration used. Where that check cannot be made, the absence is reported as a property of the available corpus and not of the language.

3. What the protocol does not require

It does not require computational methods: a proposal made by hand can satisfy all eight requirements, and many computational proposals satisfy none. It does not require giving up intuition, which is the origin of almost every good reading; it requires separating the moment of intuition from the moment of checking. And it does not require that a proposal be correct in order to be publishable: it requires that it be evaluable.

4. Why now

Three reasons. The volume of proposals has grown with the availability of tools and of prizes. The corpora are digitised, so the eight requirements are now cheap: a null is ten lines of code. And the field has produced, in recent years, examples of how it is done well: the testing against nulls of the tripartite division of Cypro-Minoan, the validation of Ithaca with an experiment among historians, and explicit warnings about the discriminating power of coincidences in small corpora. This paper only orders what those works already practise.

Declaration of assistance

The analysis of the cases and the draft of this text were produced with the assistance of a language model (Claude, Anthropic) under the author’s direction, who is responsible for all claims.

References

- Acedo, A. 2026. What can be claimed about a corpus of seven thousand signs. Draft.
- Assael, Y., T. Sommerschild, B. Shillingford et al. 2022. Restoring and attributing ancient texts using deep neural networks. *Nature* 603, 280-283.
- Corazza, M., F. Tamburini, M. Valério and S. Ferrara 2022. Unsupervised deep learning supports reclassification of Bronze Age Cypriot writing system. *PLOS ONE* 17(7), e0269544.
- Davis, B. 2025. Recent linguistics-based advances in the study of Linear A. In E. Salgarella and V. Petrakis (eds), *The Wor(l)ds of Linear A*, AURA Supplement 15, Athens, 43-56.

Duhoux, Y. 1978. Une analyse linguistique du linéaire A. In Y. Duhoux (ed.), *Études minoennes I*, Leuven, 65-129.

Godart, L. and J.-P. Olivier 1976-1985. *Recueil des inscriptions en linéaire A, I-V*. Paris.

Montecchi, B. 2009. Le frazioni, gli errori di calcolo e le unità di misura nella documentazione in lineare A. *Annali dell'Istituto Italiano di Numismatica* 55, 29-52.

Salgarella, E. 2022. Linear A. *Oxford Research Encyclopedia of Classics*.

Salgarella, E. and S. Castellan 2020. SigLA: the Signs of Linear A. <https://sigla.phis.me>

Steele, P. M. and T. Meißner 2017. From Linear B to Linear A: the problem of the backward projection of sound values. In P. M. Steele (ed.), *Understanding Relations Between Scripts*, Oxford, 93-110.

Uchitel, A. 2002-2003. HT 94. *Minos* 37-38, 81-88.